

Strategy Continuity: an intelligence layer for keeping trading knowledge inside the firm

The problem of proprietary strategies leaving with the people who built them, the documented case that proves it, what we build this weekend, and the institution it grows into.

Prepared 22 August 2026 · Under 10 pages by design · Citation codes (like [T1]) match the main Research Dossier; the cited subset is reproduced in Section 8 so this paper stands alone. Material marked **TEAM** is the team's own concept statement; material marked **BRIEF** comes from the track brief. Neither is evidence; everything else is.

1 · The concept, as stated by the team

TEAM CONCEPT STATEMENT **TEAM** (FROM TEAM CHAT, 22 AUG 2026, LIGHTLY FORMATTED)

"An intelligence layer on employee arrival across trading desks and floors, to ensure strategies unique to traders deployed at firms stay within firms upon employee departure."

"Which will in turn be used to found an institution within these firms to continue to foster growth for graduates and new employees to be in sync with the firm's current, recent and former strategies too."

"The build is the foundation, and then we do a vertical plan as part of our pitch, backed by the research."

Restated as a product thesis: quant firms pay enormous sums to acquire talent and the strategies that talent creates, but the firm's ownership of those strategies is mostly legal fiction backed by non-competes, not operational reality backed by systems. The knowledge that makes a strategy work (why parameters are set where they are, which experiments failed, which data quirks the code silently routes around) lives in the researcher's head and nowhere else. When the person leaves, the firm keeps the code and loses the strategy. We propose a research-decision capture layer that runs from an employee's first day, making the strategy's reasoning a firm-owned asset continuously rather than in a panicked exit interview, and, long term, an internal institution that uses the same captured corpus to onboard graduates into the firm's current and former strategies.

2 · Problem statement

Quantitative trading firms face a structural asymmetry: their most valuable asset is tacit knowledge, and every mechanism they use to protect it operates **after** the knowledge has already concentrated in one or two heads. The industry's actual toolkit is restrictive covenants (non-competes now reaching two years even for analysts [T4]), garden leave paid at five figures a month [T5], and litigation [T1, T6]. None of these transfer knowledge; they only delay its exit. Meanwhile churn at the pod-shop end of the industry runs at 15-20% of portfolio managers per year, with three-year retention near 55% [T8, T9], so the average strategy will outlive its author's tenure at the firm.

The problem has a mirror image that makes it worse: because strategy reasoning is undocumented, firms also cannot see changes happening inside their own models. The SEC's 2025 order against Two Sigma records an employee making unauthorized changes to more

than 14 live models over roughly two years without detection, vulnerabilities flagged internally in 2019 and unremediated until 2023, at a cost of \$90M in penalties plus \$165M voluntarily repaid [O2, T24]. Capture and change-visibility are the same missing system viewed from two sides: the firm neither knows what its people know, nor notices when what they built quietly changes.

The track brief identifies this direction ("when a quant leaves, their undocumented knowledge often leaves with them") **BRIEF**; this paper's contribution is showing the problem is documented in primary sources strong enough to build on without any interviews: a federal court filing, two SEC orders, industry turnover data, and two working precedents proving structured capture is possible and profitable.

3 · The evidence that this is real

THE NUMBERS THAT CARRY THE PITCH

- **>50% profit drop in one month after two departures**, on a strategy that made about \$1B in 2023 and was built by roughly 25 people over years: Jane Street's own amended complaint against Millennium (S.D.N.Y., 2024) [T1, T2].
- **15-20% annual PM turnover; ~55% three-year retention** at the large multi-manager funds; Millennium described as cutting PMs at 7.5% drawdowns [T8, T9].
- **\$90M penalty + \$165M repaid**: Two Sigma, for model changes nobody could see (SEC, 2025) [O2].
- **\$20M**: what Millennium paid Renaissance in 2007 over two hired quants, the same defendant seventeen years before the Jane Street case [T6].
- **Up to 2-year non-competes even for analysts** (Citadel, reported Aug 2026), and firms paying \$15-18K per month to enforce garden leave [T4, T5].
- **The industry admits it on the record**: the hedge fund lobby (MFA) formally argued to the FTC that non-competes are essential because trading IP leaves with people [T12].
- **Structured capture works and prints money**: Marshall Wace's TOPS turned structured idea submission into roughly \$30B of AUM [T14]; WorldQuant's BRAIN platform makes alpha research machine-legible by construction, with 650+ external consultants submitting through it [T13].

Three further facts shape the design. First, the knowledge really is tacit: the leading academic study of the industry, built on 300+ insider interviews, documents that trading knowledge is socially embedded and passed person to person, not through documents [T16]; Jane Street's own podcast concedes internal documentation never keeps up [F36]. Second, generic tools cannot reach it: enterprise knowledge search (Glean, \$7.2B valuation) indexes documents that exist, while quant knowledge lives in notebooks, parameter choices and failed-experiment memory that never becomes a document [T22, T23]. Third, firms already invest heavily in the adjacent, easier half: Man Group built and open-sourced ArcticDB so hundreds of quants share versioned research *data* [F20]; no equivalent exists for research *reasoning*. That asymmetry (data platforms everywhere, reasoning platforms nowhere) is the gap this concept fills.

4 · Application scenario: Jane Street, February 2024

This section grounds the product in the best-documented instance of the problem on the public record, so the pitch can reference a named firm with a court-filed account. Every factual claim below is from the amended complaint and settlement coverage [T1, T2, T3]; the counterfactual is ours.

What happened, per the court record

Jane Street built an India options strategy over years, involving roughly 25 people, generating about \$1B in profits in 2023 and described internally as "incredibly valuable" and IP-sensitive. Two traders, Douglas Schadewald and Daniel Spottiswood, resigned in February 2024 and joined Millennium. Per the complaint, the strategy's profits fell more than 50% in March 2024, the month after they left, and a new market participant appeared using the same broker with identical execution constraints. Jane Street sued, alleging trade-secret misappropriation; the defense argued the traders took only their "experience and expertise." The case settled in December 2024 on undisclosed terms [T2]. A regulator later put the franchise's scale on the record: SEBI's 2025 order describes about \$4.3B of India profits from January 2023 to March 2025 [T3].

Why this is the perfect reference case

The lawsuit's central ambiguity is precisely the product's reason to exist. Jane Street had the code, the infrastructure, and 23 of the 25 people, yet lost half the strategy's output in a month, and then could not cleanly demonstrate in court where "firm IP" ended and the traders' personal "experience and expertise" began. A firm running a continuity layer would be in a different position on both counts: operationally (the reasoning, market intuitions, and execution constraints around the strategy would exist as a structured, queryable corpus that the remaining 23 people could work from) and legally (a timestamped record of firm-hosted contributions is close to the strongest trade-secret documentation a plaintiff can have; courts weigh whether the firm took reasonable measures to protect the secret, and a capture system is such a measure by construction).

The counterfactual, stated carefully for the pitch

We do not claim the tool would have prevented the departures, and we do not claim the profits would have been untouched; skill is real and some knowledge cannot be serialized. The claim is narrower and defensible: the firm's recovery time and its legal position both improve with a decision-genealogy corpus, and the gap between "we kept the code" and "we kept the strategy" narrows. The complaint itself documents which knowledge mattered by fighting over it: models, signals, execution constraints, broker arrangements. Those are exactly the artifacts the layer captures.

The buyer-side scenario: Millennium as the high-churn profile

The same case shows the other customer. Millennium runs 300+ autonomous pods with 15-20% PM turnover and has now twice been the defendant in landmark quant knowledge disputes (2007 Renaissance, \$20M; 2024 Jane Street) [T6, T1, T8]. A firm whose operating model is deliberately high-churn is structurally dependent on knowledge surviving personnel change, and currently has no system that makes it do so. Two Sigma is the third scenario: its \$255M outcome was not about departure at all but about the absence of a record of who changed which model, when, and why [O2, T24]; the same capture layer is that record.

5 · What we build this weekend

Working name: **Continuity** (rename freely). The hackathon prototype is the intelligence layer's core loop, demonstrated on a realistic synthetic trading-research repo (strategy code, notebooks, parameter files, a fake chat log) that we author on Friday so the demo is credible without any proprietary data.

Prototype scope (the foundation)

- **Ambient capture.** Git hooks and a Jupyter extension that intercept every merged change to strategy code or parameters and require a one-line structured rationale (what changed, why, what was tried and rejected). Each entry auto-links the diff, author, timestamp, and affected strategy. This is the decision genealogy: the record whose absence cost Two Sigma \$255M [O2] and whose presence would have been Jane Street's evidence [T1].
- **The debrief agent.** An LLM interviewer that runs short scheduled sessions with each researcher ("you changed the vol filter threshold twice this month; walk me through why"), generating targeted questions from the capture stream rather than generic prompts, and filing transcripts against the right strategy. Modeled on the proven structured-capture precedents [T13, T14] but aimed at reasoning, not ideas.
- **Knowledge concentration metrics.** A bus-factor dashboard per strategy: who is the sole author or toucher of which components, weighted by P&L attribution; a departure simulation showing what fraction of a strategy's attributed code and decisions would lose their last knowledgeable owner.
- **The handover pack.** One command generates, for any (simulated) departing employee, a structured succession document: their strategies, their decision history, the open questions only they can answer ranked by undocumentedness, and a ready-made interview script for their final weeks.

Demo script (three minutes)

Open on the Jane Street numbers: a \$1B strategy, two resignations, profits down more than 50% in a month, per the firm's own court filing [T1, T2]. Show the synthetic firm's dashboard: one strategy glowing red (bus factor 1). Simulate that person resigning. Generate the handover pack live; show the debrief agent conducting the exit session and the knowledge map turning from red toward green as answers file in. Close on the two-sided value: continuity for the desk, and a timestamped IP record the general counsel would have paid for in April 2024.

Why the research says this is the right build

SOURCE FINDING	DESIGN DECISION IT FORCED
Knowledge is tacit and never becomes documents [T16, F36]; enterprise search only indexes what exists [T22]	Capture at the moment of action (commits, parameter changes), not via wikis or after-the-fact documentation drives
Exit interviews and wikis demonstrably fail for tacit knowledge [T23]	Continuous ambient capture from day one TEAM , with the exit pack as output, not input
Structured capture succeeded when it was lightweight and incentive-aligned (TOPS, BRAIN) [T14, T13]	One-line rationales and short agent sessions, never long forms; visible credit for captured contributions
The Two Sigma failure was invisibility of change, not absence of talent [O2]	The same stream doubles as a model change-control audit trail, which also answers "why would a firm buy this"
Firms already accept research-infrastructure mandates (ArcticDB for data) [F20]	Position as research infrastructure, not surveillance; it sits beside the data platform as the reasoning platform

6 · Long-term vision and vertical plan

The long-term thesis, per the team concept **TEAM**: the capture layer's corpus becomes the foundation of an institution inside the firm, so that continuity stops being an emergency process and becomes how the firm teaches itself.

Phase 1 (hackathon to seed): the continuity layer.

The prototype above, hardened: on-prem or VPC-only deployment (no quant firm ships this data to a SaaS cloud), integration with git, Jupyter, and chat, and the bus-factor and change-control dashboards. Wedge buyer: the COO and general counsel, because the same artifact serves succession risk and trade-secret documentation. The reference sale writes itself from the litigation record cited above.

Phase 2 (the institution): onboarding built on the corpus.

Once decisions and debriefs accumulate, the corpus becomes courseware. New graduates onboard against the firm's actual current, recent, and former strategies **TEAM**: guided walkthroughs of a live strategy's decision history, "why is this parameter 0.7" answered by the recorded rationale rather than by interrupting the one person who knows. This attacks the documented onboarding gap (sink-or-swim is the industry's stated norm; documentation never keeps up [F36]) and shortens the time-to-productive that firms currently buy with two-year non-competes on the other end.

Phase 3 (the flywheel): garden leave and alumni knowledge.

Firms pay departing quants five figures monthly to sit idle for up to two years [T4, T5]. A compliance-safe, firm-side structured debrief program during that paid window converts dead spend into curriculum, and the corpus becomes the firm's institutional memory of strategies past: searchable, teachable, and legally owned. At this point the "institution within the firm" exists in fact: capture feeds curriculum, curriculum feeds new researchers, whose work is captured in turn.

Why now. Three trends converge: non-compete law is under active attack (the FTC's attempted ban and state-level movement [T12]), which weakens the legal substitute for actual knowledge retention; LLMs make the debrief agent and corpus querying feasible for the first time; and regulators are converging on the same artifact from the compliance side (model change-control expectations post-Two Sigma [O2], algo inventories and documented change management under MiFID II RTS 6 [O10, O11]). The product's record is simultaneously a retention asset, a legal asset, and a compliance asset.

7 · Objections we will face, answered honestly

- **"Traders will refuse to feed it."** True if it reads as surveillance or as making them replaceable. Mitigations built into the design: capture is about strategies, not people-monitoring; contribution is visible and creditable (the TOPS lesson: people feed systems that pay them credit [T14]); friction is kept to one line. We say this openly in the pitch; pretending incentives away would be the weak move.
- **"Tacit knowledge cannot be fully captured."** Correct, and the pitch never claims total capture. The claim is marginal: the difference between profits falling 50% and falling meaningfully less, and between an unprovable trade secret and a documented one. The court record shows the marginal value is measured in hundreds of millions [T1, T2].
- **"Firms would build this in-house."** The strongest objection, since firms build everything [F18-F22]. Two answers: they have not (the confessed tooling for data exists, for reasoning does not), and the compliance-artifact framing gives an external product a neutral-record advantage an internal tool lacks in litigation.
- **"Is this just a wiki with hooks?"** No: wikis wait for voluntary prose after the fact and demonstrably fail [T23]. This captures at the point of action, interrogates actively, and measures concentration. The deliverable is a genealogy and a risk metric, not pages.

8 · Pitch reference card and citations

Every claim in the pitch or video should be spoken with its source. The quotable core, then the cited subset in full (codes match the main Research Dossier).

SAY THIS	BACKED BY
"A \$1B-a-year strategy lost over half its profits in the month after two people resigned. That is Jane Street's own number, in its own federal court filing."	Amended complaint, S.D.N.Y. 2024 [T1]; settlement coverage [T2]
"The multi-manager funds turn over 15 to 20 percent of their portfolio managers every year. The average strategy will outlive its author's seat."	Net Interest, "Peak Pod" [T8]; Hedgeweek/WSJ [T9]
"Two Sigma paid \$90M and repaid \$165M because nobody could see who changed 14 live models over two years. The missing system is the same one."	SEC press release 2025-15 [O2]; MoFo alert [T24]
"The industry's own lobby told the FTC that non-competes are essential because strategies leave with people. That is the industry admitting it has no operational answer, only a legal one."	FTC record and MFA opposition [T12]
"Structured capture is proven: Marshall Wace turned it into roughly \$30B of AUM, and WorldQuant industrialized it into a platform."	Net Interest, "Alpha Capture" [T14]; WorldQuant BRAIN [T13]
"Firms built world-class platforms for research data. Nobody has built one for research reasoning."	ArcticDB and peers [F20, F18, F19]; gap analysis, Dossier 2.7

Cited sources

- T1** · Jane Street amended complaint v. Millennium, Schadewald and Spottiswood, S.D.N.Y. 1:24-cv-02783, Doc. 64 (2024). Strategy built by ~25 people; profits down >50% the month after two departures; trade-secret vs "experience and expertise" dispute. https://www.sewkis.com/wp-content/uploads/Jane-Street_amended-complaint.pdf
- T2** · Settlement coverage (Bloomberg-sourced), Dec 2024; ~\$1B 2023 profits on the disputed strategy. <https://www.claimsjournal.com/news/national/2024/12/09/327780.htm>
- T3** · SEBI interim order coverage, July 2025: ~\$4.3B India profits Jan 2023-Mar 2025. Primary PDF retrievable from [sebi.gov.in](https://blogs.law.ox.ac.uk/oblb/blog-post/2025/07/jane-street-and-expiry-day-trap-unpacking-sebis-crackdown-algorithmic). <https://blogs.law.ox.ac.uk/oblb/blog-post/2025/07/jane-street-and-expiry-day-trap-unpacking-sebis-crackdown-algorithmic>
- T4** · Bloomberg, Aug 2026: Citadel's 1-2 year non-competes down to analyst level (paywalled; corroboration linked). <https://www.briefs.co/news/citadel-s-two-year-ban-on-job-moves-stalls-analyst-careers/>
- T5** · Institutional Investor, 2022: \$15-18K/month garden-leave compliance payments, Citadel v. Prieur. <https://www.institutionalinvestor.com/article/b1zyzmvhxf2612/An-Employee-Left-Citadel-Securities-for-a-Startup-Now-He-s-Facing-a-Lawsuit>
- T6** · Renaissance v. Millennium: \$20M settlement over two hired quants (2007), confirmed by counsel. <https://www.cravath.com/renaissance-technologies-corp-settles-trade-secrets-dispute/>
- T8** · Rubinstein, "Peak Pod," Net Interest: 15-20% annual PM turnover; ~55% three-year retention; Millennium ~300 teams. <https://www.netinterest.co/p/peak-pod>
- T9** · Hedgeweek (WSJ-derived), 2025: nine-figure packages; PMs cut at 7.5% drawdowns. <https://www.hedgeweek.com/multi-manager-hedge-funds-escalate-talent-war-arms-race-with-nine-figure-pay-packages/>
- T12** · FTC non-compete rule (2024) and MFA opposition: the industry's on-record argument that trading IP leaves with people. <https://www.ftc.gov/news-events/news/press-releases/2024/04/ftc-announces-rule-banning-noncompetes>
- T13** · WorldQuant BRAIN and Global Alphathon: machine-legible alpha research; 650+ consultants. <https://www.worldquant.com/brain/>
- T14** · Rubinstein, "Alpha Capture," Net Interest: Marshall Wace TOPS, ~500,000 structured ideas/year by 2005, ~\$30B AUM. <https://www.netinterest.co/p/alpha-capture>
- T16** · MacKenzie, "Trading at the Speed of Light" (Princeton, 2021): 300+ insider interviews; trading knowledge is tacit and socially embedded. <https://press.princeton.edu/books/hardcover/9780691211381/trading-at-the-speed-of-light>
- T22** · Glean (enterprise search, \$7.2B valuation) indexes existing documents; quant reasoning never becomes one. https://en.wikipedia.org/wiki/Glean_Technologies

- T23** · HR literature on wikis and exit interviews failing to capture tacit knowledge. <https://hrexecutive.com/unshared-knowledge-destroys-a-companys-bottom-line/>
- T24** · Morrison Foerster on the Two Sigma order: no systematic record of who changed which model parameter, when, and why. <https://www.mofo.com/resources/insights/250226-top-5-sec-enforcement-developments>
- O2** · SEC press release 2025-15, Two Sigma: unauthorized changes to 14+ live models; \$90M penalty, \$165M repaid; flagged 2019, fixed 2023. <https://www.sec.gov/newsroom/press-releases/2025-15>
- O10** · FCA multi-firm review, 2018: firms could not inventory their own algorithms; limited dev/test/deploy documentation. <https://www.fca.org.uk/publication/multi-firm-reviews/algorithmic-trading-compliance-wholesale-markets.pdf>
- O11** · Central Bank of Ireland thematic review, 2021: RTS 6 self-assessments deficient; documentation of controls lacking. <https://www.centralbank.ie/docs/default-source/regulation/industry-market-sectors/investment-firms/mifid-firms/regulatory-requirements-and-guidance/thematic-assessment-of-algorithmic-trading-firms-compliance.pdf>
- F18** · HRT, distributed filesystem for research at hundreds of petabytes. <https://www.hudsonrivertrading.com/hrtbeat/distributed-filesystem-for-scalable-research/>
- F19** · Two Sigma, "Treating Data as Code." <https://www.twosigma.com/articles/treating-data-as-code-at-two-sigma/>
- F20** · Man Group, "Our Journey Creating ArcticDB": the research data platform firms do build; no reasoning equivalent exists. <https://arcticdb.io/blog/our-journey-creating-arcticdb-solving-the-challenge-of-dataframes-at-scale/>
- F22** · G-Research, Armada: research compute platform, open-sourced. <https://armadaproject.io/>
- F36** · Jane Street, Signals and Threads, "Writing, Technically": internal documentation never keeps up. <https://signalsandthreads.com/writing-technically/>

Integrity notes for the video: the >50% figure is Jane Street's allegation in a settled case; say "per Jane Street's court filing," not "proven in court." Millennium's 7.5% drawdown rule and turnover figures are industry reporting, not filings. The Two Sigma numbers are from the SEC's own press release and are safe to state plainly.